

# Relational Algebra

CS 355 Database Management System Design

S. Taneja  
Oct 1, 2019




Saw this somewhere on the Internet yesterday....

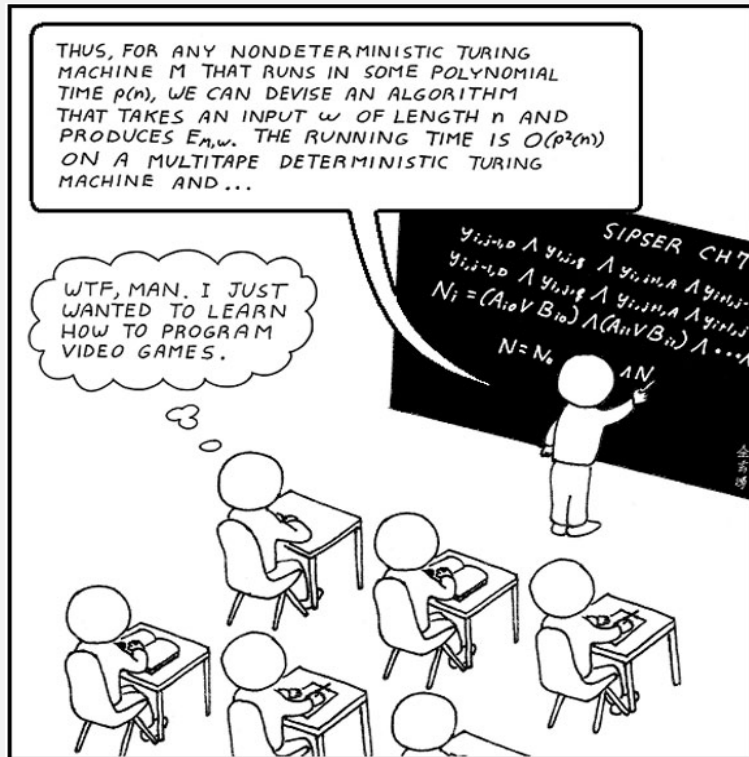
# Most Popular Programming Languages

Source: Stack Overflow Developer Survey 2019


WOMEN WHO  
**CODE**®



## Course Content



[http://abstrusegoose.com/strips/computer\\_science\\_major.PNG](http://abstrusegoose.com/strips/computer_science_major.PNG)



*"If you find that you're spending almost all your time on theory, start turning some attention to practical things; it will improve your theories. If you find that you're spending almost all your time on practice, start turning some attention to theoretical things; it will improve your practice."* **Don Knuth**

## Formal Relational Languages

**Relational algebra** and **relational calculus** are two formal languages defined for the relational model. In contrast, **SQL** was the practical language for the relational model.

Both the algebra and the calculus were defined before SQL and in many ways form the basis of SQL and its implementation.

**Relational algebra** defines the basic **operators** for the relational model and specifies the rules for their use and composition.

**IMP:** Relational algebra *operators* are applied to *relations* and produce *relations* as their result.

**Queries**, or retrieval requests, can be specified as **relational algebra expressions**.

## Some more facts about Relational Algebra

Relational algebra is inherently **procedural** and forms the basis of how SQL statements are implemented and executed. *(tell us HOW to do)*

**Relational calculus** is a higher-level **declarative** language for specifying relational queries. *(tells us WHAT to do) – we will not cover this in CS355*

Both relational algebra and relational calculus were defined by E.F. Codd, with elements of the algebra being defined in the landmark original paper on the relational model. **Codd's Theorem** states that the algebra and the calculus are logically equivalent.

# Relational algebra

**Basic Operators**      Required to provide full expressive power

**Projection**

$$\rho_{attributes}(R)$$

Project attribute values from all tuples of R

**Selection**

$$S_{condition}(R)$$

Select tuples from R that satisfy the condition

**Cross-Product**

$$R \times S$$

Concatenate each tuple in R with each tuple in S

**Union**

$$R \cup S$$

All tuples in either R or S

**Difference**

$$R - S$$

The tuples in R that are not in S

**Rename**

$$Name \leftarrow Expression$$

Names the result of an expression or renames an attribute

# Relational algebra

**Additional Operators**   Convenient but not strictly necessary

**Intersection**

$$R \cap S$$

Tuples in both R and S

**Theta Join**

$$R \bowtie_C S$$

Combined tuples from R and S that satisfy the condition C

**Equijoin**

$$R \bowtie_C S$$

Special case of theta join where C uses only equality

**Natural Join**

$$R \bowtie S$$

Join based on equality of shared attributes

**Left Outer Join**

$$R \bowtie_{\theta} S$$

Join with non-matching tuples from R

**Right Outer Join**

$$R \bowtie_{\theta} S$$

Join with non-matching tuples from S

**Full Outer Join**

$$R \bowtie_{\theta} S$$

Join with non-matching tuples from R and S

**Division**

$$R(XY) \div S(Y)$$

Tuples in R that match with all relevant tuples in S

# Relational algebra

**Extended Operators**    Add missing functionality

## **Aggregate Functions**

Specify mathematical functions (e.g., counting, summing, etc.) on values or groups of values.

$$\langle grouping \rangle \mathcal{F} \langle function \rangle (R)$$



# Relational algebra interpreter

RA

<http://www.cs.duke.edu/~junyang/ra/>

RA supports the following relational algebra operators:

`\select_{cond}` is the relational selection operator. For example, to select Drinker tuples with name Amy or Ben, we can write `\select_{name = 'Amy' or name = 'Ben'} Drinker;`. Syntax for `cond` follows SQL. Note that string literals should be enclosed in single quotes, and you may use boolean operators `and`, `or`, and `not`. Comparison operators `<=`, `<`, `=`, `>`, `>=`, and `<>` work on both string and numeric types. For string match you can use the SQL LIKE operator; e.g., `\select_{name like 'A%'} drinker;` finds all drinkers whose name start with A, as % is a wildcard character that matches any number of characters.

`\project_{attr_list}` is the relational projection operator, where `attr_list` is a comma-separated list of attribute names. For example, to find out what beers are served by Talk of the Town (but without the price information), we can write `\project_{bar, beer} (\select_{bar = 'Talk of the Town'} Serves);`

`\join_{cond}` is the relational theta-join operator. For example, to join Drinker(name, address) and Frequents(drinker, bar, times\_a\_week) relations together using drinker name, we can write `Drinker \join_{name = drinker} Frequents;`. Syntax for `cond` again follows SQL; see notes on `\select` for more details.

`\join` is the relational natural join operator. For example, to join Drinker(name, address) and Frequents(drinker, bar, times\_a\_week) relations together using drinker name, we can write `Drinker \join \rename_{name, bar, times_a_week} Frequents;`. Natural join will automatically equate all pairs of identically named attributes from its inputs (in this case, name), and output only one attribute per pair. Here we use `\rename` to create two name attributes for the natural join; see notes on `\rename` below for more details.

`\cross` is the relational cross product operator. For example, to compute the cross product of Drinker and Frequents, we can write `Drinker \cross Frequents;`

`\union`, `\diff`, and `\intersect` are the relational union, difference, and intersect operators. For a trivial example, to compute the union, difference, and intersection between Drinker and itself, we can write `Drinker \union Drinker;`, `Drinker \diff Drinker;`, and `Drinker \intersect Drinker;`, which would return Drinker itself, an empty relation, and Drinker itself, respectively.

`\rename_{new_attr_name_list}` is the relational rename operator, where `new_attr_name_list` is a comma-separated list of new names, one for each attribute of the input relation. For example, to rename the attributes of relation Drinker and compute the cross product of Drinker and itself, we can write `\rename_{name1, address1} Drinker \cross \rename_{name2, address2} Drinker;`

## Sample relations for example queries *Based on Codd's "suppliers" database*

### Suppliers

snum	sname	status	city
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens
S6	Brown	15	Berlin

### Parts

pnum	pname	color	weight	city
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London
P7	Gear	Yellow	18	Berlin

### Projects

jnum	jname	city
J1	Sorter	Paris
J2	Punch	Rome
J3	Reader	Athens
J4	Console	Athens
J5	Filler	London
J6	Layer	Oslo
J7	Tape	London

### Shipments

snum	pnum	jnum	qty
S1	P1	J1	200
S1	P1	J4	700
S2	P3	J1	400
S2	P3	J2	200
S2	P3	J3	200
S2	P3	J4	500
S2	P3	J5	600
S2	P3	J6	400
S2	P3	J7	800
S2	P5	J2	100
S3	P3	J1	200
S3	P4	J2	500
S4	P6	J3	300
S4	P6	J7	300
S5	P2	J2	200
S5	P2	J4	100
S5	P5	J5	500
S5	P5	J7	100
S1	P4	J1	100
S1	P6	J2	200

## Projection

 $\rho_{attributes}(R)$ 

Projects the values of the specified attributes from all tuples in R.

 $\rho_{snum}(Suppliers)$ 

```
ra> \project_{snum}(Suppliers);
```

Output schema: (snum text)

-----

S1

S2

S3

S4

S5

S6

-----

Total number of rows: 6

Suppliers

snum	sname	status	city
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens
S6	Brown	15	Berlin

The result of the PROJECT operation can be visualized as a **vertical partition** of the relation into two relations:

- One with the desired columns (attributes)
- The other contains the discarded columns

## Projection

$\rho_{attributes}(R)$

Projects the values of the specified attributes from all tuples in R.

$\rho_{pnum}(Shipments)$

```
ra> \project_{pnum}(Shipments);
```

Output schema: (pnum text)

-----

P1

P2

P3

P4

P5

P6

-----

Total number of rows: 6

Shipments

snum	pnum	jnum	qty
S1	P1	J1	200
S1	P1	J4	700
S2	P3	J1	400
S2	P3	J2	200
S2	P3	J3	200
S2	P3	J4	500
S2	P3	J5	600
S2	P3	J6	400
S2	P3	J7	800
S2	P5	J2	100
S3	P3	J1	200
S3	P4	J2	500
S4	P6	J3	300
S4	P6	J7	300
S5	P2	J2	200
S5	P2	J4	100
S5	P5	J5	500
S5	P5	J7	100
S1	P4	J1	100
S1	P6	J2	200

## Projection

$\rho_{city}(Suppliers)$

```
ra> \project_{city}(Suppliers);
```

Output schema: (city text)

-----

Athens

Berlin

London

Paris

-----

Total number of rows: 4

$\rho_{jnum,jname}(Projects)$

```
ra> \project_{jnum, jname}(Projects);
```

Output schema: (jnum text, jname text)

-----

J1|Sorter

J2|Punch

J3|Reader

J4|Console

J5|Filler

J6|Layer

J7|Tape

-----

Total number of rows: 7

### Observations:

1. The result of a PROJECT operation has only the attributes specified in <attribute list> in the same order as they appear in the list. Thus, its **degree** is equal to the number of attributes in <attribute list>
2. \*If the attribute list included only nonkey attributes of R, duplicate tuples are likely to occur (see e.g. on the left above). The PROJECT operation *removes any duplicate tuples*.

## Projection

$\rho_{pnum, jnum}(Shipments)$

```
ra> \project_{pnum, jnum} (Shipments);
```

Output schema: (pnum text, jnum text)

```
-----  
P1|J1  
P1|J4  
P2|J2  
P2|J4  
P3|J1  
P3|J2  
P3|J3  
P3|J4  
P3|J5  
P3|J6  
P3|J7  
P4|J1  
P4|J2  
P5|J2  
P5|J5  
P5|J7  
P6|J2  
P6|J3  
P6|J7  
-----
```

Total number of rows: 19

Shipments

snum	pnum	jnum	qty
S1	P1	J1	200
S1	P1	J4	700
S2	P3	J1	400
S2	P3	J2	200
S2	P3	J3	200
S2	P3	J4	500
S2	P3	J5	600
S2	P3	J6	400
S2	P3	J7	800
S2	P5	J2	100
S3	P3	J1	200
S3	P4	J2	500
S4	P6	J3	300
S4	P6	J7	300
S5	P2	J2	200
S5	P2	J4	100
S5	P5	J5	500
S5	P5	J7	100
S1	P4	J1	100
S1	P6	J2	200

## Selection

$S_{condition}(R)$

Selects tuples from R that satisfy the given condition.

- Think of it as a **horizontal partition of a relation** into two set of tuples
- one that satisfies the condition and are selected,
  - Others that do not satisfy and are filtered out

$S_{status>15}(Suppliers)$

```
ra> \select_{status > 15}(Suppliers);
```

-----

S1|Smith|20|London

S3|Blake|30|Paris

S4|Clark|20|London

S5|Adams|30|Athens

-----

Total number of rows: 4

Suppliers

snum	sname	status	city
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens
S6	Brown	15	Berlin

## Selection

Note that the SELECT operator is unary; that is, it is applied to a single relation.

Moreover, SELECT is applied to each tuple individually.

The **degree** of the relation resulting from a SELECT operation – its number of attributes – is same as degree of R.

SELECT operation is commutative, that is

$$\sigma_{\langle \text{cond1} \rangle}(\sigma_{\langle \text{cond2} \rangle}(R)) = \sigma_{\langle \text{cond2} \rangle}(\sigma_{\langle \text{cond1} \rangle}(R))$$



## Selection

$S_{qty < 200}$  (*Shipments*)

```
ra> \select_{qty < 200}(Shipments);
```

-----

S1|P4|J1|100

S2|P5|J2|100

S5|P2|J4|100

S5|P5|J7|100

-----

Total number of rows: 4

Shipments

snum	pnum	jnum	qty
S1	P1	J1	200
S1	P1	J4	700
S2	P3	J1	400
S2	P3	J2	200
S2	P3	J3	200
S2	P3	J4	500
S2	P3	J5	600
S2	P3	J6	400
S2	P3	J7	800
S2	P5	J2	100
S3	P3	J1	200
S3	P4	J2	500
S4	P6	J3	300
S4	P6	J7	300
S5	P2	J2	200
S5	P2	J4	100
S5	P5	J5	500
S5	P5	J7	100
S1	P4	J1	100
S1	P6	J2	200

## Selection

$S_{color='Blue' OR city='London'}(Parts)$

```
ra> \select_{color = 'Blue' OR city = 'London'}(Parts);
```

-----

```
P1|Nut|Red|12.0|London  
P3|Screw|Blue|17.0|Rome  
P4|Screw|Red|14.0|London  
P5|Cam|Blue|12.0|Paris  
P6|Cog|Red|19.0|London
```

-----

```
Total number of rows: 5
```

## Operator composition

$\rho_{snum}(\mathcal{S}_{status>15}(Suppliers))$

```
ra> \project_{snum}(\select_{status > 15}(Suppliers));
```

-----

S1

S3

S4

S5

-----

Total number of rows: 4

Suppliers

snum	sname	status	city
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens
S6	Brown	15	Berlin

## Operator composition

$\rho_{pnum, pname}(S_{color='Blue' OR city='London'}(Parts))$

```
ra> \project_{pnum, pname}(\select_{color = 'Blue' OR city = 'London'}(Parts));
```

```
-----
```

```
P1|Nut
```

```
P3|Screw
```

```
P4|Screw
```

```
P5|Cam
```

```
P6|Cog
```

```
-----
```

```
Total number of rows: 5
```

## Cross product

$R \text{ ' } S$

Concatenates all tuples in R with all tuples in S

**R1**

A	B	C
1	2	3
4	2	1
5	6	7

**R2**

D	E	A
8	2	4
1	5	5
9	1	4

$R1 \text{ ' } R2$

```
ra> R1 \cross R2;
```

Output schema: (A integer, B integer, C integer, D integer, E integer, A:1 integer)

-----

```
1|2|3|1|5|5
1|2|3|8|2|4
1|2|3|9|1|4
4|2|1|1|5|5
4|2|1|8|2|4
4|2|1|9|1|4
5|6|7|1|5|5
5|6|7|8|2|4
5|6|7|9|1|4
```

-----

Total number of rows: 9

## Cross product

### *Suppliers ´ Parts*

```
ra> Suppliers \cross Parts;
```

Output schema: (snum text, sname text, status integer, city text, pnum text, pname text, color text, weight float, city:1 text)

-----

```
S1|Smith|20|London|P1|Nut|Red|12.0|London
S1|Smith|20|London|P2|Bolt|Green|17.0|Paris
S1|Smith|20|London|P3|Screw|Blue|17.0|Rome
S1|Smith|20|London|P4|Screw|Red|14.0|London
S1|Smith|20|London|P5|Cam|Blue|12.0|Paris
S1|Smith|20|London|P6|Cog|Red|19.0|London
S1|Smith|20|London|P7|Gear|Yellow|18.0|Berlin
S2|Jones|10|Paris|P1|Nut|Red|12.0|London
S2|Jones|10|Paris|P2|Bolt|Green|17.0|Paris
S2|Jones|10|Paris|P3|Screw|Blue|17.0|Rome
S2|Jones|10|Paris|P4|Screw|Red|14.0|London
S2|Jones|10|Paris|P5|Cam|Blue|12.0|Paris
```

. . .

```
S6|Brown|15|Berlin|P5|Cam|Blue|12.0|Paris
S6|Brown|15|Berlin|P6|Cog|Red|19.0|London
S6|Brown|15|Berlin|P7|Gear|Yellow|18.0|Berlin
```

-----

Total number of rows: 42

## Cross product

### *Suppliers ´ Shipments*

```
ra> Suppliers \cross Shipments;
```

Output schema: (snum text, sname text, status integer, city text, snum:1 text, pnum text, jnum text, qty integer)

-----

```
S1|Smith|20|London|S1|P1|J1|200
S1|Smith|20|London|S1|P1|J4|700
S1|Smith|20|London|S1|P4|J1|100
```

. . .

```
S4|Clark|20|London|S2|P3|J7|800
S4|Clark|20|London|S2|P5|J2|100
S4|Clark|20|London|S3|P3|J1|200
S6|Brown|15|Berlin|S5|P2|J2|200
S6|Brown|15|Berlin|S5|P2|J4|100
S6|Brown|15|Berlin|S5|P5|J5|500
S6|Brown|15|Berlin|S5|P5|J7|100
```

-----

Total number of rows: 120

## Cross product

$S_{snum=snum} (Suppliers \times Shipments)$

```
\select_{snum = snum}(Suppliers \cross Shipments); // Not in RA  
// must rename duplicate attributes; messy but ...
```

-----

```
S1|Smith|20|London|S1|P1|J1|200  
S1|Smith|20|London|S1|P1|J4|700  
S1|Smith|20|London|S1|P4|J1|100  
S1|Smith|20|London|S1|P6|J2|200  
S2|Jones|10|Paris|S2|P3|J1|400  
S2|Jones|10|Paris|S2|P3|J2|200  
S2|Jones|10|Paris|S2|P3|J3|200  
S2|Jones|10|Paris|S2|P3|J4|500  
S2|Jones|10|Paris|S2|P3|J5|600  
S2|Jones|10|Paris|S2|P3|J6|400  
S2|Jones|10|Paris|S2|P3|J7|800  
S2|Jones|10|Paris|S2|P5|J2|100  
S3|Blake|30|Paris|S3|P3|J1|200  
S3|Blake|30|Paris|S3|P4|J2|500  
S4|Clark|20|London|S4|P6|J3|300  
S4|Clark|20|London|S4|P6|J7|300  
S5|Adams|30|Athens|S5|P2|J2|200  
S5|Adams|30|Athens|S5|P2|J4|100  
S5|Adams|30|Athens|S5|P5|J5|500  
S5|Adams|30|Athens|S5|P5|J7|100
```

-----

Total number of rows: 20



## Union

$R \cup S$

The tuples in either R or S. R and S must be “union compatible.”

### Undergrad

sid	sname	gpa
9021	Pat	4.00
9022	Kelly	3.85
9023	Logan	3.25

### Grad

sid	sname	gpa
9024	Taylor	3.89
9025	Jessie	3.37

$Undergrad \cup Grad$

```
ra> Undergrad \union Grad;
Output schema: (sid text, sname text, GPA float)
-----
9021|Pat|4.0
9022|Kelly|3.85
9023|Logan|3.25
9024|Taylor|3.89
9025|Jessie|3.37
-----
Total number of rows: 5
```

## Union

$$\pi_{pnum}(\sigma_{city='Paris'}(Parts)) \cup \pi_{pnum}(\sigma_{jnum='J2'}(Shipments))$$

This returns the

```
ra>
\project_{pnum}(\select_{city = 'Paris'}(Parts))
  \union
\project_{pnum}(\select_{jnum = 'J2'}(Shipments));
```

Output schema: (pnum text)

-----

P2

P3

P4

P5

P6

-----

Total number of rows: 5

## Difference

$R - S$

The tuples in R that are not in S. R and S must be “union compatible.”

**Undergrad**

sid	sname	gpa
9021	Pat	4.00
9022	Kelly	3.85
9023	Logan	3.25

**Rotc**

sid	sname	gpa
9021	Pat	4.00

$Undergrad - Rotc$

```
ra> Undergrad \diff Rotc;
```

```
Output schema: (sid text, sname text, GPA
float)
```

```
-----
```

```
9022|Kelly|3.85
```

```
9023|Logan|3.25
```

```
-----
```

```
Total number of rows: 2
```

## Difference

$$\rho_{pnum}(Parts) - \rho_{pnum}(S_{jnum='J2'}(Shipments))$$

This returns



```
ra> \project_{pnum}(Parts) \diff \project_{pnum}(\select_{jnum = 'J2'}(Shipments));
```

Output schema: (pnum text)

-----

P1

P7

-----

Total number of rows: 2

## RENAME Operation

- In general, for most queries, we need to apply **several** relational algebra operations one after the other.

$$\pi_{\text{Fname, Lname, Salary}}(\sigma_{\text{Dno}=5}(\text{EMPLOYEE}))$$

- Either we can write the operations as a single **relational algebra expression** by nesting the operations, or we can apply one operation at a time and create intermediate result relations.
  - In the latter case, we must give names to the relations that hold the intermediate results.

$$\begin{aligned}\text{DEP5\_EMPS} &\leftarrow \sigma_{\text{Dno}=5}(\text{EMPLOYEE}) \\ \text{RESULT} &\leftarrow \pi_{\text{Fname, Lname, Salary}}(\text{DEP5\_EMPS})\end{aligned}$$

- The general **RENAME** operation when applied to a relation  $R$  of degree  $n$  is denoted by any of the following three forms:

$$\rho_{S(B_1, B_2, \dots, B_n)}(R) \quad \text{or} \quad \rho_S(R) \quad \text{or} \quad \rho_{(B_1, B_2, \dots, B_n)}(R)$$

symbol  $\rho$  (rho) = RENAME operator

$S$  = new relation name

$B_1, B_2, \dots, B_n$  are the new attribute names.

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name and id number of all suppliers in Paris having status greater than 20.

$$\rho_{sname, snum} (S_{(city='Paris') \text{ and } (status > 20)} (Suppliers))$$

```
ra> \project_{sname, snum}(\select_{city = 'Paris' AND status > 20}(Suppliers));
```

Output schema: (sname text, snum text)

-----

Blake|S3

-----

Total number of rows: 1

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name of suppliers who supply part P2.

$$\rho_{sname} (S_{pnum='P2'} (S_{snum=snum} (Suppliers \Join Shipments)))$$

```
ra>
\project_{sname}(
  \select_{pnum = 'P2'}(
    \select_{snum = sid}(
      Suppliers
      \cross
      \rename_{sid, pnum, jnum, qty} Shipments
    )
  )
);
```

Output schema: (sname text)


-----  
Adams  
-----

Total number of rows: 1

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name of suppliers who supply part P2.



Since the cross product almost always combines too much, we will typically follow it with a selection. The **join** operator is simply the composition of a product and a selection.



# Relational algebra

**Additional Operators**   Convenient but not strictly necessary

**Intersection**

$$R \cap S$$

Tuples in both R and S

**Theta Join**

$$R \bowtie_C S$$

Combined tuples from R and S that satisfy the condition C

**Equijoin**

$$R \bowtie_C S$$

Special case of theta join where C uses only equality

**Natural Join**

$$R \bowtie S$$

Join based on equality of shared attributes

**Left Outer Join**

$$R \Join_{\theta} S$$

Join with non-matching tuples from R

**Right Outer Join**

$$R \Join_{\theta} S$$

Join with non-matching tuples from S

**Full Outer Join**

$$R \Join_{\theta} S$$

Join with non-matching tuples from R and S

**Division**

$$R(XY) \div S(Y)$$

Tuples in R that match with all relevant tuples in S

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name of suppliers who supply part P2.

$$\pi_{sname}(\sigma_{pnum='P2'}(Suppliers \bowtie Shipments))$$

```
ra>
\project_{sname}(\select_{pnum = 'P2'}(Suppliers \join Shipments));
```

Output schema: (sname text)

-----

Adams

-----

Total number of rows: 1

## Three “inner” joins

**Theta Join**      $R \bowtie_C S$      Combined tuples from R and S that satisfy the condition C

**Equijoin**      $R \bowtie_C S$      Special case of theta join where condition uses only equality

**Natural Join**      $R \bowtie S$      Join based on equality of shared attributes

## Natural join

**Natural Join**  $R \bowtie S$  Join based on equality of **all** shared attributes

Shared attributes appear only once in resulting schema

R	A	B	C
	1	2	3
	4	2	1
	5	6	7

S	D	E	A	C
	8	2	4	3
	1	5	5	2
	9	1	4	1

$R \bowtie S$	A	B	C	D	E
	4	2	1	9	1

```
ra> R \join S;
```

```
Output schema: (A integer, B integer, C integer, D integer, E integer)
```

```
-----  
4|2|1|9|1  
-----
```

```
Total number of rows: 1
```

## Natural join

*Suppliers database*

**Query:** Retrieve the name of suppliers who supply part P2.

$$\pi_{sname}(\sigma_{pnum='P2'}(Suppliers \bowtie Shipments))$$

*The natural join is appropriate since snum is the only shared attribute between Suppliers and Shipments.*

```
ra> \project_{sname}(Suppliers \join (\select_{pnum = 'P2'}(Shipments)));
```

Output schema: (sname text)

-----

Adams

-----

Total number of rows: 1

## Natural join

*Suppliers database*

**Query:** Retrieve the name of suppliers who supply a blue part.

$$\pi_{sname}(\sigma_{color='Blue'}((Parts \bowtie Shipments) \bowtie Suppliers))$$

```
\project_{sname}(\select_{color = 'Blue'}(Parts \join Shipments \join Suppliers));
```

Output schema: (sname text)

-----

Jones

-----

Total number of rows: 1

**Incorrect! Adams and Blake are missing.**

*The natural join is incorrect since city becomes one of the shared attributes.*

# Equijoin

Equijoin

$R \bowtie_C S$

Join based on equality comparisons only

Shared attributes appear twice in resulting schema

R	A	B	C
	1	2	3
	4	2	1
	5	6	7

S	D	E	A	C
	8	2	4	3
	1	5	5	2
	9	1	4	1

A	B	C	D	E	A	C
4	2	1	8	2	4	3
4	2	1	9	1	4	1
5	6	7	1	5	5	2

$R \bowtie_{R.A=S.A} S$

```
ra> R \join_{A = SA} \rename_{D,E,SA,SC} S;
```

```
Output schema: (A integer, B integer, C integer, D integer, E  
integer, SA integer, SC integer)
```

```
-----
```

```
4|2|1|8|2|4|3
```

```
4|2|1|9|1|4|1
```

```
5|6|7|1|5|5|2
```

```
-----
```

```
Total number of rows: 3
```

# Equijoin

*Suppliers database*

**Query:** Retrieve the name of suppliers who supply a blue part.

$$\pi_{sname}(\sigma_{color='Blue'}((Parts \bowtie Shipments) \bowtie_{snum=snum} Suppliers))$$

```
ra> \project_{sname}(\select_{color = 'Blue'}(Parts \join Shipments \join_{snum =
sid} \rename_{sid,sname,status,city} Suppliers));
```

Output schema: (sname text)

-----

Adams

Blake

Jones

-----

Total number of rows: 3

*The equijoin is correct since it allows us to specify which shared attribute to use.*



# Theta join

Theta Join

$R \bowtie_C S$

Join condition based on any boolean condition.

All attributes of both R and S are in the resulting schema.

R	A	B	C
	1	2	3
	4	2	1
	5	6	7

S	D	E	A	C
	8	2	4	3
	1	5	5	2
	9	1	4	1

A	B	C	D	E	A	C
5	6	7	8	2	4	3
5	6	7	9	1	4	1

$R \bowtie_{R.A > S.A} S$

```
ra> R \join_{A > SA} \rename_{D,E,SA,SC} S;
```

Output schema: (A integer, B integer, C integer, D integer, E integer, SA integer, SC integer)

```
-----  
5|6|7|8|2|4|3  
5|6|7|9|1|4|1  
-----
```

Total number of rows: 2

## Theta join

*Suppliers database*

**Query:** Retrieve the supplier numbers of all suppliers with a status greater than the status of S1.

$$\pi_{snum}(\sigma_{snum='S1'}(Suppliers) \bowtie_{status_{right} > status_{left}} (Suppliers))$$

```
\project_{snum} ( \rename_{stat1}(\project_{status}(\select_{snum =  
'S1'}(Suppliers))) \join_{status > stat1} Suppliers );
```

Output schema: (snum text)

-----

S3

S5

-----

Total number of rows: 2

## Algebraic properties of joins

Inner joins are **associative** and **commutative** operators.

R	A	B	C
	1	2	3
	4	2	1
	5	6	7

S	D	E	A	C
	8	2	4	3
	1	5	5	2
	9	1	4	1

T	F	A
	2	3
	3	4
	4	5

**Associative**

$$R \bowtie S \bowtie T = (R \bowtie S) \bowtie T = R \bowtie (S \bowtie T)$$

```
ra> (R \join S) \join T;
```

Output schema: (A integer, B integer, C integer, D integer, E integer, F integer)

-----

```
4|2|1|9|1|3
```

-----

Total number of rows: 1

```
ra> R \join (S \join T);
```

Output schema: (A integer, B integer, C integer, D integer, E integer, F integer)

-----

```
4|2|1|9|1|3
```

-----

Total number of rows: 1

## Algebraic properties of joins

Inner joins are **associative** and **commutative** operators.

R	A	B	C
	1	2	3
	4	2	1
	5	6	7

S	D	E	A	C
	8	2	4	3
	1	5	5	2
	9	1	4	1

**Commutative**

$$R \bowtie S = S \bowtie R$$

```
ra> R \join S;
```

Output schema: (A integer, B integer, C integer, D integer, E integer)

```
-----  
4|2|1|9|1
```

```
-----  
Total number of rows: 1
```

```
ra> S \join R;
```

Output schema: (D integer, E integer, A integer, C integer, B integer)

```
-----  
9|1|4|1|2
```

```
-----  
Total number of rows: 1
```

## Queries in relational algebra

### *Suppliers database*

**Query:** Retrieve the number, name, and city of all projects. If any project is located in the same city as a supplier, list the supplier number as well.

$$\pi_{jnum, jname, city, snum} (Projects \bowtie_{Projects.city = Suppliers.city} Suppliers)$$

```
ra> \project_{jnum,jname,city1,snum}( (\rename_{jnum,jname,city1}(Projects))
\join_{city1 = city2} (\rename_{snum,sname,status,city2}(Suppliers)) );
```

Output schema: (jnum text, jname text, city1 text, snum text)

-----

```
J1|Sorter|Paris|S2
J1|Sorter|Paris|S3
J3|Reader|Athens|S5
J4|Console|Athens|S5
J5|Filler|London|S1
J5|Filler|London|S4
J7|Tape|London|S1
J7|Tape|London|S4
```

-----

Total number of rows: 8

**Incorrect – J2 and J6 are missing.**

## Three “outer” joins

Left Outer Join  $R \bowtie_{\theta} S$

Join with non-matching tuples from R

Right Outer Join  $R \bowtie_{\theta} S$

Join with non-matching tuples from S

Full Outer Join  $R \bowtie_{\theta} S$

Join with non-matching tuples from R and S

R	A	B	C
	1	2	3
	4	2	1
	5	6	7

S	F	A
	2	3
	3	4
	4	5

$R \bowtie S$

A	B	C	F	A
1	2	3	NULL	NULL
4	2	1	3	4
5	6	7	4	5

$R \bowtie S$

A	B	C	F	A
NULL	NULL	NULL	2	3
4	2	1	3	4
5	6	7	4	5

## Queries in relational algebra

**Query:** Retrieve the number, name, and city of all projects along with the supplier numbers, if any, that are located in the same city as a project.

$\pi_{jnum, jname, city, snum} (Projects \bowtie_{Projects.city = Suppliers.city} Suppliers)$

RA doesn't support outer joins, so this output is from PostgreSQL:

jnum	jname	city	snum
J1	Sorter	Paris	S3
J1	Sorter	Paris	S2
J2	Punch	Rome	
J3	Reader	Athens	S5
J4	Console	Athens	S5
J5	Filler	London	S4
J5	Filler	London	S1
J6	Layer	Oslo	
J7	Tape	London	S4
J7	Tape	London	S1

(10 rows)

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name and number of all parts that are carried in any of the following colors: red, yellow, green.

$$\pi_{pname,plum}(\sigma_{color='Red' \text{ OR } color='Yellow' \text{ OR } color='Green'}(Parts))$$

```
ra> \project_{pname, pnum}(\select_{color = 'Red' OR color = 'Yellow' OR color = 'Green'}(Parts));
```

Output schema: (pname text, pnum text)

```
-----  
Bolt|P2  
Cog|P6  
Gear|P7  
Nut|P1  
Screw|P4  
-----
```

Total number of rows: 5



## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name and number of all parts that are carried in any of the following colors: red, yellow, green.

$$\pi_{pname,plum}(\sigma_{color='Red'}(Parts) \cup \sigma_{color='Yellow'}(Parts) \cup \sigma_{color='Green'}(Parts))$$

```
ra>
\project_{pname, pnum}(
  \select_{color = 'Red'}(Parts)
  \union
  \select_{color = 'Yellow'}(Parts)
  \union
  \select_{color = 'Green'}(Parts)
);
```

Output schema: (pname text, pnum text)

-----

Bolt|P2

Cog|P6

Gear|P7

Nut|P1

Screw|P4

-----

Total number of rows: 5

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name and number of all suppliers who supply at least one red part.

$$\pi_{sname, snum} (Suppliers \bowtie Shipments \bowtie (\pi_{pnum} (\sigma_{color='Red'} (Parts))))$$

```
ra> \project_{sname, snum}(Suppliers \join Shipments \join
(\project_{pnum}(\select_{color = 'Red'}(Parts))));
```

Output schema: (sname text, snum text)

```
-----
Blake|S3
Clark|S4
Smith|S1
```

```
-----
Total number of rows: 3
```

## Queries in relational algebra

### *Suppliers database*

**Query:** Retrieve the name and number of all suppliers who supply either a red part or a blue part.

$$\pi_{sname, snum}(Suppliers \bowtie Shipments \bowtie (\pi_{pnum}(\sigma_{color='Red' \text{ OR } color='Blue'}(Parts))));$$

```
ra> \project_{sname, snum}(Suppliers \join Shipments \join
(\project_{pnum}(\select_{color = 'Red' OR color = 'Blue'}(Parts))));
```

Output schema: (sname text, snum text)

-----

Adams|S5

Blake|S3

Clark|S4

Jones|S2

Smith|S1

-----

Total number of rows: 5

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name and number of all suppliers who supply both a red part and a blue part.

$$\pi_{sname, snum}(Suppliers \bowtie Shipments \bowtie (\pi_{pnum}(\sigma_{color='Red' \text{ AND } color='Blue'}(Parts))));$$

**A part can only have one color.**

```
ra> \project_{sname, snum}(Suppliers \join Shipments \join
(\project_{pnum}(\select_{color = 'Red' AND color = 'Blue'}(Parts))));
```

Output schema: (sname null, snum null)

-----

-----

Total number of rows: 0

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name and number of all suppliers who supply both a red part and a blue part.

$$\pi_{sname, snum} (Suppliers \bowtie Shipments \bowtie (\pi_{pnum} (\sigma_{color='Red'} (Parts)))) \\ \cap \pi_{sname, snum} (Suppliers \bowtie Shipments \bowtie (\pi_{pnum} (\sigma_{color='Blue'} (Parts))));$$

```
ra> \project_{sname, snum}(
    Suppliers \join Shipments
    \join (\project_{pnum}(\select_{color = 'Red'}(Parts))))

\intersect

\project_{sname, snum}(
    Suppliers \join Shipments
    \join (\project_{pnum}(\select_{color = 'Blue'}(Parts))));
```

Output schema: (sname text, snum text)

-----

Blake|S3

-----

Total number of rows: 1

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name, number, and color of all parts that are supplied by supplier S3.

$$\pi_{pname, pnum, color} (Parts \bowtie (\pi_{pnum} (\sigma_{snum='S3'} (Shipments))));$$

```
ra> \project_{pname, pnum, color}(Parts \join
(\project_{pnum}(\select_{snum='S3'}(Shipments))));
```

Output schema: (pname text, pnum text, color text)

-----

Screw|P3|Blue

Screw|P4|Red

-----

Total number of rows: 2

What if we need the details of the parts which are not supplied by S3? Should we change = to != ?

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the number of all parts that are not supplied by supplier S3.

$$\pi_{pnum}(Parts) - \pi_{pnum}(\sigma_{snum='S3'}(Shipments))$$

```
ra> \project_{pnum}(Parts) \diff \project_{pnum}(\select_{snum = 'S3'}(Shipments));
```

Output schema: (pnum text)

-----

P1

P2

P5

P6

P7

-----

Total number of rows: 5

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name and number of suppliers who are currently not shipping any parts.

$$\pi_{sname, snum}(Suppliers) - \pi_{sname, snum}(Suppliers \bowtie Shipments)$$

```
ra> \project_{sname, snum}(Suppliers) \diff \project_{sname, snum}(Suppliers
\join Shipments);
```

Output schema: (sname text, snum text)

-----

Brown|S6

-----

Total number of rows: 1

When using difference, be sure that the schemas on both sides of the '-' are similar in structure.



# Queries in relational algebra

Suppliers database

This is a “*for all*” query, and the division operator is specifically intended to be used for these type queries.

**Query:** Retrieve the name and number of suppliers who are currently shipping all the parts that are carried.

$A(XY) \div B(Y)$  Returns the X values in A that are paired with ALL the Y values in B.

A(snum, pnum) = suppliers and the parts they ship

B(pnum) = all parts carried

A	
snum	pnum
S1	P1
S1	P2
S1	P3
S1	P4
S2	P1
S2	P2
S3	P2
S4	P2
S4	P4

B1	
pnum	
P2	

A ÷ B1	
snum	
S1	
S2	
S3	
S4	

B2	
pnum	
P2	
P4	

A ÷ B2	
snum	
S1	
S4	

B3	
pnum	
P1	
P2	
P4	

A ÷ B3	
snum	
S1	

## Queries in relational algebra

### *Suppliers database*

**Query:** Retrieve the name and number of suppliers who are currently shipping all the parts that are carried.

This is a “for all” query, and the division operator is specifically intended to be used for these type queries.

**Division**  $A(XY) \div B(Y)$  The X values that are paired in A with all the Y values in B.

*Suppliers and the parts they ship*

snum	pnum
------	------

*All parts carried*

pnum
------

*Suppliers shipping all parts*

snum
------

$$\rho_{snum, pnum}(Shipments) \div \rho_{pnum}(Parts) =$$

The relation returned by division operator will return those tuples from relation A which are associated with every B's tuple.

# Queries in relational algebra

## *Suppliers database*

**Query:** Retrieve the id number of suppliers who are currently shipping all the parts that are carried.

This is a “for all” query, and the division operator is specifically intended to be used for these type queries.

**Division**  $A(XY) \div B(Y)$

The X values that are paired in A with all the Y values in B.

*Suppliers and the parts they ship*

snum	pnum
S1	P1
S2	P3
S2	P5
S3	P3
S3	P4
S4	P6
S5	P2
S5	P5

÷

*All parts carried*

pnum
P1
P2
P3
P4
P5
P6
P7

=

*Suppliers shipping all parts*

snum
------

*None for the example data.*

# Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the id number of suppliers who are currently shipping to all the projects.

$$\pi_{snum, jnum}(Shipments) \div \pi_{jnum}(Projects)$$

*Suppliers and the  
projects they ship to*

snum	jnum
S1	J1
S1	J2
S1	J4
S2	J1
S2	J2
S2	J3
S2	J4
S2	J5
S2	J6
S2	J7
S3	J1
S3	J2
S4	J3
S4	J7
S5	J2
S5	J4
S5	J5
S5	J7

*All  
projects*

jnum
J1
J2
J3
J4
J5
J6
J7

÷

=

*Suppliers shipping  
to all projects*

snum
S2

## Division in terms of core operations

Just like multiplication and division are inverse operators in arithmetic, cross product and division are inverse operators in relational algebra.

$$C(XY) \leftarrow A(X) \times B(Y)$$

$$A(X) = C(XY) \div B(Y)$$

$$B(Y) = C(XY) \div A(X)$$

## Division in terms of core operations

Given **A(XY)** and **B(Y)**:  $A \div B = \pi_X(A) - \pi_X((\pi_X(A) \times B) - A)$

$A \leftarrow \pi_{snum,jnum}(Shipments)$

$B \leftarrow \pi_{jnum}(Projects)$

$A \div B = \pi_{snum}(A) - \pi_{snum}((\pi_{snum}(A) \times B) - A)$

## Division in terms of core operations

Given **A(XY)** and **B(Y)**:  $A \div B = \pi_X(A) - \pi_X((\pi_X(A) \times B) - A)$

$A \leftarrow \pi_{snum,jnum}(Shipments)$

$B \leftarrow \pi_{jnum}(Projects)$

$A \div B = \pi_{snum}(A) - \pi_{snum}(\underbrace{(\pi_{snum}(A) \times B)}_{\text{All possible pairings of snum and jnum}} - A)$

All possible pairings of  
snum and jnum

## Division in terms of core operations

Given **A(XY)** and **B(Y)**:  $A \div B = \pi_X(A) - \pi_X((\pi_X(A) \times B) - A)$

$A \leftarrow \pi_{snum,jnum}(Shipments)$

$B \leftarrow \pi_{jnum}(Projects)$

$A \div B = \pi_{snum}(A) - \pi_{snum}(\underbrace{(\pi_{snum}(A) \times B) - A}_{\text{The pairings of snum and jnum that do not appear in our data}})$

The pairings of snum and jnum  
that do not appear in our data



## Division in terms of core operations

Given  $A(XY)$  and  $B(Y)$ :  $A \div B = \pi_X(A) - \pi_X((\pi_X(A) \times B) - A)$

$A \leftarrow \pi_{snum,jnum}(Shipments)$

$B \leftarrow \pi_{jnum}(Projects)$

$A \div B = \pi_{snum}(A) - \pi_{snum}((\pi_{snum}(A) \times B) - A)$

The suppliers that are not the ones we're  
looking for

## Division in terms of core operations

Given **A(XY)** and **B(Y)**:  $A \div B = \pi_X(A) - \pi_X((\pi_X(A) \times B) - A)$

$A \leftarrow \pi_{snum,jnum}(Shipments)$

$B \leftarrow \pi_{jnum}(Projects)$

$A \div B = \pi_{snum}(A) - \pi_{snum}((\pi_{snum}(A) \times B) - A)$

The suppliers that are left in our data set after we remove the ones we're not looking for

## Division in terms of core operations

Given **A(XY)** and **B(Y)**:  $A \div B = \pi_X(A) - \pi_X((\pi_X(A) \times B) - A)$

$A \leftarrow \pi_{snum,jnum}(Shipments)$

$B \leftarrow \pi_{jnum}(Projects)$

$A \div B = \pi_{snum}(A) - \pi_{snum}((\pi_{snum}(A) \times B) - A)$

```
ra> \project_{snum}(\project_{snum,jnum}(Shipments)) \diff \project_{snum}( (
(\project_{snum}(\project_{snum,jnum}(Shipments))) \cross
\project_{jnum}(Projects) ) \diff \project_{snum,jnum}(Shipments) );
```

Output schema: (snum text)

-----

S2

-----

Total number of rows: 1

## Queries in relational algebra

*Suppliers database*

**Query:** Retrieve the name and number of suppliers who are currently shipping all the red parts that are being shipped.

$$SuppliersShippingParts \leftarrow \pi_{snum, pnum}(Shipments)$$
$$RedPartsBeingShipped \leftarrow \pi_{pnum}((\sigma_{color='Red'}(Parts)) \bowtie Shipments)$$
$$ShippingAllRed \leftarrow SuppliersShippingParts \div RedPartsBeingShipped$$
$$\pi_{sname, snum}(Suppliers \bowtie ShippingAllRed)$$

Now try this on radb