

SQL

CS 355 Database Management System Design

S. Taneja



SQL

SQL (Structured Query Language) is a language designed to manage and query data stored in a relational database. SQL is based on relational tuple calculus and some algebra.

Originally named SEQUEL (Structured English Query Language), SQL was developed by **Chamberlin and Royce** at IBM, as part of the **System R project**.

The first commercial version of SQL was made available in 1979 by Oracle, followed shortly by IBM.

SQL was adopted as an ANSI standard in 1986, as an ISO standard in 1987. The standard has undergone many revisions and updates, with the **latest being SQL:2016 (ISO/IEC 9075:2016)**.

SQL is supported by a large number of database products including IBM's **DB2 (commercial version of system R)**, ORACLE, SYBASE, SQLServer, MySQL, and many more. I'm unaware of any major database vendor that fully supports the SQL standard. Many SQL implementations are, at least in some respects, vendor-dependent.

Sample relations for example queries *Based on Codd's "suppliers" database*

Suppliers

snum	sname	status	city
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens
S6	Brown	15	Berlin

Parts

pnum	pname	color	weight	city
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London
P7	Gear	Yellow	18	Berlin

Projects

jnum	jname	city
J1	Sorter	Paris
J2	Punch	Rome
J3	Reader	Athens
J4	Console	Athens
J5	Filler	London
J6	Layer	Oslo
J7	Tape	London

Shipments

snum	pnum	jnum	qty
S1	P1	J1	200
S1	P1	J4	700
S2	P3	J1	400
S2	P3	J2	200
S2	P3	J3	200
S2	P3	J4	500
S2	P3	J5	600
S2	P3	J6	400
S2	P3	J7	800
S2	P5	J2	100
S3	P3	J1	200
S3	P4	J2	500
S4	P6	J3	300
S4	P6	J7	300
S5	P2	J2	200
S5	P2	J4	100
S5	P5	J5	500
S5	P5	J7	100
S1	P4	J1	100
S1	P6	J2	200

Why Filter Data?

Be specific about the data that we want to retrieve

Reduce the number of records we retrieve

Increase query performance

Reduce the strain on the client application

And so forth...

SQL SELECT statement

Most basic form:

```
SELECT <attribute list>  
FROM   <table list>
```

Basic form:

```
SELECT <attribute list>  
FROM   <table list>  
WHERE  <condition>;
```

More complete form:

```
SELECT <attribute list>  
FROM   <table list>  
[WHERE <condition>]  
[GROUP BY <grouping attribute(s)>  
[HAVING <group condition>]]  
[ORDER BY <attribute list>]
```

Basic form, core relational algebra operators, and joins

SQL SELECT statement

Most basic form:

```
SELECT <attribute list>  
FROM   <table list>
```

```
SELECT snum  
FROM   suppliers;
```

```
snum  
-----  
S1  
S2  
S3  
S4  
S5  
S6  
(6 rows)
```

SQL SELECT statement

Most basic form:

```
SELECT <attribute list>  
FROM   <table list>
```

```
SELECT snum, sname  
FROM   suppliers;
```

snum		sname
S1		Smith
S2		Jones
S3		Blake
S4		Clark
S5		Adams
S6		Brown

(6 rows)

SQL SELECT statement

Most basic form:

```
SELECT <attribute list>  
FROM   <table list>
```

```
SELECT snum  
FROM   shipments;
```

A list of rows:

```
snum  
-----  
S1  
S1  
S2  
S2  
S2  
S2  
S2  
S2  
S2  
S2  
S2  
S3  
S3  
S4  
S4  
S5  
S5  
S5  
S5  
S1  
S1  
(20 rows)
```

SQL SELECT statement

Most basic form:

```
SELECT <attribute list>  
FROM   <table list>
```

```
SELECT DISTINCT snum  
FROM   shipments;
```

A list of rows:

```
snum  
-----  
S3  
S1  
S4  
S5  
S2  
(5 rows)
```

SQL SELECT statement

Basic/common form:

```
SELECT <attribute list>
FROM   <table list>
WHERE  <condition>;
```

```
SELECT snum, sname
FROM   suppliers
WHERE  city = 'Paris';
```

snum		sname
S2		Jones
S3		Blake

(2 rows)

SQL SELECT statement

Basic/common form:

```
SELECT sname, pnum, qty
FROM   suppliers, shipments;
```

Cross product, not a join

sname	pnum	qty
Smith	P1	200
Smith	P1	700
Smith	P3	400
Smith	P3	200
Smith	P3	200
Smith	P3	500
Smith	P3	600
Smith	P3	400
Smith	P3	800
Smith	P5	100
Smith	P3	200
. . .		

. . .		
Brown	P3	800
Brown	P5	100
Brown	P3	200
Brown	P4	500
Brown	P6	300
Brown	P6	300
Brown	P2	200
Brown	P2	100
Brown	P5	500
Brown	P5	100
Brown	P4	100
Brown	P6	200
(120 rows)		

SQL SELECT statement

Basic/common form:

```
SELECT sname, pnum, qty
FROM   suppliers, shipments
WHERE  suppliers.snum = shipments.snum;
```

sname	pnum	qty
Smith	P1	200
Smith	P1	700
Jones	P3	400
Jones	P3	200
Jones	P3	200
Jones	P3	500
Jones	P3	600
Jones	P3	400
Jones	P3	800
Jones	P5	100

Blake	P3	200
Blake	P4	500
Clark	P6	300
Clark	P6	300
Adams	P2	200
Adams	P2	100
Adams	P5	500
Adams	P5	100
Smith	P4	100
Smith	P6	200

(20 rows)

Quick digression – order of evaluation

SQL is more declarative than procedural, but there is an order of evaluation that is usually observed.

Different DBMS implementations could behave differently, but this is typical.

```
5 ---SELECT <attribute list>  
1 ---FROM   <table list>  
2 ---[WHERE <condition>]  
3 ---[GROUP BY <grouping attribute(s)>  
4 ---[HAVING <group condition>]]  
6 ---[ORDER BY <attribute list>]
```

SQL – joins

```
SELECT sname, pnum, qty  
FROM   suppliers, shipments  
WHERE  suppliers.snum = shipments.snum;
```

This is semantically equivalent to a join, but SQL-92 explicitly added join syntax.

“ANSI joins” are specified as the following table expressions in the FROM clause.

(inner joins)

```
table1 JOIN table2 ON condition
```

```
table1 JOIN table2 USING (columns)
```

```
table1 NATURAL JOIN table2
```

There’s also a (redundant) CROSS JOIN operator (“unqualified join”):

```
table1 CROSS JOIN table2
```

SQL – joins

```
SELECT sname, pnum, qty
FROM   suppliers JOIN shipments ON suppliers.snum = shipments.snum;
```

sname	pnum	qty
Smith	P1	200
Smith	P1	700
Jones	P3	400
Jones	P3	200
Jones	P3	200
Jones	P3	500
Jones	P3	600
Jones	P3	400
Jones	P3	800
Jones	P5	100

Blake	P3	200
Blake	P4	500
Clark	P6	300
Clark	P6	300
Adams	P2	200
Adams	P2	100
Adams	P5	500
Adams	P5	100
Smith	P4	100
Smith	P6	200

(20 rows)

SQL – joins

```
SELECT sname, pnum, qty
FROM   suppliers JOIN shipments USING (snum);
```

sname	pnum	qty
Smith	P1	200
Smith	P1	700
Jones	P3	400
Jones	P3	200
Jones	P3	200
Jones	P3	500
Jones	P3	600
Jones	P3	400
Jones	P3	800
Jones	P5	100

Blake	P3	200
Blake	P4	500
Clark	P6	300
Clark	P6	300
Adams	P2	200
Adams	P2	100
Adams	P5	500
Adams	P5	100
Smith	P4	100
Smith	P6	200

(20 rows)

SQL – joins

```
SELECT sname, pnum, qty
FROM   suppliers NATURAL JOIN shipments;
```

sname	pnum	qty
-----+	-----+	-----
Smith	P1	200
Smith	P1	700
Jones	P3	400
Jones	P3	200
Jones	P3	200
Jones	P3	500
Jones	P3	600
Jones	P3	400
Jones	P3	800
Jones	P5	100

Blake	P3	200
Blake	P4	500
Clark	P6	300
Clark	P6	300
Adams	P2	200
Adams	P2	100
Adams	P5	500
Adams	P5	100
Smith	P4	100
Smith	P6	200

(20 rows)

SQL – joins

Outer joins are also provided in SQL by using the keywords LEFT, RIGHT, or FULL in conjunction with any of the three inner join forms (ON, USING, NATURAL).

```
SELECT snum, pnum, suppliers.city  
FROM   suppliers JOIN parts USING (city);
```

Inner join

snum	pnum	city
S1	P6	London
S1	P4	London
S1	P1	London
S2	P5	Paris
S2	P2	Paris
S3	P5	Paris
S3	P2	Paris
S4	P6	London
S4	P4	London
S4	P1	London
S6	P7	Berlin

(11 rows)

SQL – joins

Outer joins are also provided in SQL by using the keywords LEFT, RIGHT, or FULL in conjunction with any of the three inner join forms (ON, USING, NATURAL).

```
SELECT snum, pnum, suppliers.city  
FROM   suppliers LEFT JOIN parts USING (city);
```

Left outer join

snum	pnum	city
S1	P6	London
S1	P4	London
S1	P1	London
S2	P5	Paris
S2	P2	Paris
S3	P5	Paris
S3	P2	Paris
S4	P6	London
S4	P4	London
S4	P1	London
S5		Athens
S6	P7	Berlin

(12 rows)

SQL – joins

Outer joins are also provided in SQL by using the keywords LEFT, RIGHT, or FULL in conjunction with any of the three inner join forms (ON, USING, NATURAL).

```
SELECT snum, pnum, suppliers.city  
FROM   suppliers RIGHT JOIN parts USING (city);
```

Right outer join

snum	pnum	city
S1	P6	London
S1	P4	London
S1	P1	London
S2	P5	Paris
S2	P2	Paris
S3	P5	Paris
S3	P2	Paris
S4	P6	London
S4	P4	London
S4	P1	London
S6	P7	Berlin
	P3	

(12 rows)

SQL – joins

Outer joins are also provided in SQL by using the keywords LEFT, RIGHT, or FULL in conjunction with any of the three inner join forms (ON, USING, NATURAL).

```
SELECT snum, pnum, suppliers.city
FROM   suppliers FULL JOIN parts USING (city);
```

Full outer join

snum	pnum	city
S1	P6	London
S1	P4	London
S1	P1	London
S2	P5	Paris
S2	P2	Paris
S3	P5	Paris
S3	P2	Paris
S4	P6	London
S4	P4	London
S4	P1	London
S5		Athens
S6	P7	Berlin
	P3	

(13 rows)

SQL – union and difference

SQL supports all the core relational algebra operators: projection, selection, cross product, union, and difference.

```
SELECT pnum FROM parts WHERE city = 'Paris'  
UNION  
SELECT pnum FROM shipments WHERE jnum = 'J2';
```

pnum
P3
P6
P2
P4
P5

(5 rows)

*The union operation **eliminates** duplicates before returning the result.*

SQL – union and difference

SQL supports all the core relational algebra operators: projection, selection, cross product, union, and difference.

```
SELECT pnum FROM parts WHERE city = 'Paris'  
UNION ALL  
SELECT pnum FROM shipments WHERE jnum = 'J2';
```

pnum

P2
P5
P3
P5
P4
P2
P6
(7 rows)

The union all operation leaves duplicates in the result.

SQL – union and difference

SQL supports all the core relational algebra operators: projection, selection, cross product, union, and difference.

```
SELECT pnum FROM parts
```

EXCEPT

```
SELECT pnum FROM shipments WHERE jnum = 'J2';
```

```
pnum  
-----  
P1  
P7  
(2 rows)
```

*The except operation eliminates duplicates before returning the result.
If duplicates are needed, EXCEPT ALL should be used.*

SQL – intersection

```
SELECT pnum FROM parts WHERE city = 'Paris'  
INTERSECT  
SELECT pnum FROM shipments WHERE jnum = 'J2';
```

```
pnum  
-----  
P5  
P2  
(2 rows)
```

The intersect operation eliminates duplicates before returning the result. If duplicates are needed, INTERSECT ALL should be used.

Example queries using these operations

Queries in SQL

Suppliers database

Query: Retrieve the name and id number of all suppliers in Paris having status greater than 20.

```
SELECT sname, snum
FROM   suppliers
WHERE  city = 'Paris'
AND    status > 20;
```

sname		snum
-----+		-----
Blake		S3

(1 row)

Queries in SQL

Suppliers database

Query: Retrieve the name and id number of all suppliers in Paris having status greater than 20.

```
(SELECT sname, snum
FROM   suppliers
WHERE  city = 'Paris')
INTERSECT
(SELECT sname, snum
FROM   suppliers
WHERE  status > 20);
```

sname		snum
-----	+	-----
Blake		S3

(1 row)

Queries in SQL

Suppliers database

Query: Retrieve the name of suppliers who supply part P2.

```
SELECT sname
FROM   suppliers, shipments
WHERE  shipments.pnum = 'P2';
```

Incorrect

```
sname
-----
Smith
Smith
Jones
Jones
Blake
Blake
Clark
Clark
Adams
Adams
Brown
Brown
(12 rows)
```

Queries in SQL

Suppliers database

Query: Retrieve the name of suppliers who supply part P2.

```
SELECT sname
FROM   suppliers, shipments
WHERE  suppliers.snum = shipments.snum
AND    shipments.pnum = 'P2';
```

```
sname
-----
Adams
Adams
(2 rows)
```

Queries in SQL

Suppliers database

Query: Retrieve the name of suppliers who supply part P2.

```
SELECT sname
FROM   suppliers NATURAL JOIN shipments
WHERE  shipments.pnum = 'P2';
```

```
sname
-----
Adams
Adams
(2 rows)
```


Queries in SQL

Suppliers database

Query: Retrieve the name of suppliers who supply part P2.

```
SELECT DISTINCT sname
FROM   suppliers, shipments
WHERE  suppliers.snum = shipments.snum
AND    shipments.pnum = 'P2';
```

```
sname
-----
Adams
(1 row)
```

```
SELECT DISTINCT sname
FROM   suppliers NATURAL JOIN shipments
WHERE  shipments.pnum = 'P2';
```

```
SELECT DISTINCT sname
FROM   suppliers JOIN shipments USING (snum)
WHERE  shipments.pnum = 'P2';
```

```
SELECT DISTINCT sname
FROM   suppliers JOIN shipments ON (suppliers.snum = shipments.snum)
WHERE  shipments.pnum = 'P2';
```

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all parts that are carried in any of the following colors: red, yellow, green.

```
SELECT pname, pnum
FROM   parts
WHERE  color = 'Red' OR color = 'Yellow' OR color = 'Green';
```

pname		pnum
Nut		P1
Bolt		P2
Screw		P4
Cog		P6
Gear		P7

(5 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all parts that are carried in any of the following colors: red, yellow, green.

```
(SELECT pname, pnum FROM parts WHERE color = 'Red')  
UNION  
(SELECT pname, pnum FROM parts WHERE color = 'Yellow')  
UNION  
(SELECT pname, pnum FROM parts where color = 'Green');
```

pname		pnum
Nut		P1
Gear		P7
Bolt		P2
Cog		P6
Screw		P4

(5 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all suppliers who supply at least one red part.

```
SELECT suppliers.sname, suppliers.snum
FROM   suppliers, shipments, parts
WHERE  suppliers.snum = shipments.snum
AND    parts.pnum = shipments.pnum
AND    parts.color = 'Red';
```

sname		snum
Smith		S1
Smith		S1
Blake		S3
Clark		S4
Clark		S4
Smith		S1
Smith		S1

(7 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all suppliers who supply at least one red part.

```
SELECT DISTINCT suppliers.sname, suppliers.snum
FROM   suppliers, shipments, parts
WHERE  suppliers.snum = shipments.snum
AND    parts.pnum = shipments.pnum
AND    parts.color = 'Red';
```

sname		snum
Clark		S4
Blake		S3
Smith		S1

(3 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all suppliers who supply at least one red part.

```
SELECT DISTINCT suppliers.sname, suppliers.snum
FROM   suppliers NATURAL JOIN shipments NATURAL JOIN parts
WHERE  parts.color = 'Red';
```



Remember what a
natural join does!

sname		snum
Clark		S4
Smith		S1

(2 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all suppliers who supply at least one red part.

```
SELECT DISTINCT suppliers.sname, suppliers.snum
FROM   suppliers, shipments, parts
WHERE  suppliers.snum = shipments.snum
AND    parts.pnum = shipments.pnum
AND    parts.color = 'Red';
```

```
SELECT DISTINCT suppliers.sname, suppliers.snum
FROM   suppliers NATURAL JOIN shipments
       JOIN parts USING (pnum)
WHERE  parts.color = 'Red';
```

sname	snum
Clark	S4
Blake	S3
Smith	S1

(3 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all suppliers who supply at least one red part. List the results in ascending order of last name.

```
SELECT DISTINCT suppliers.sname, suppliers.snum
FROM   suppliers NATURAL JOIN shipments
      JOIN parts USING (pnum)
WHERE  parts.color = 'Red'
ORDER BY suppliers.sname;
```

sname		snum
Blake		S3
Clark		S4
Smith		S1

(3 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all suppliers who supply either a red part or a blue part.

```
SELECT DISTINCT suppliers.sname, suppliers.snum
FROM    suppliers
        JOIN shipments USING (snum)
        JOIN parts USING (pnum)
WHERE   (parts.color = 'Red' OR parts.color = 'Blue');
```

sname		snum
-----+		-----
Jones		S2
Blake		S3
Smith		S1
Clark		S4
Adams		S5

(5 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all suppliers who supply either a red part or a blue part.

```
(SELECT sname, snum
FROM   parts
      JOIN shipments USING (pnum)
      JOIN suppliers USING (snum)
WHERE  color = 'Red')
UNION
(SELECT sname, snum
FROM   parts
      JOIN shipments USING (pnum)
      JOIN suppliers USING (snum)
WHERE  color = 'Blue');
```

sname	snum
Blake	S3
Jones	S2
Smith	S1
Clark	S4
Adams	S5

(5 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all suppliers who supply both a red part and a blue part.

```
(SELECT sname, snum
FROM    parts JOIN shipments USING (pnum)
        JOIN suppliers USING (snum)
WHERE   color = 'Red')
INTERSECT
(SELECT sname, snum
FROM    parts JOIN shipments USING (pnum)
        JOIN suppliers USING (snum)
WHERE   color = 'Blue');
```

sname	snum
Blake	S3

(1 row)

Queries in SQL

Suppliers database

Query: Retrieve all pairs of suppliers that are located in the same city.

```
SELECT suppliersA.snum, suppliersB.snum
FROM   suppliers AS suppliersA, suppliers AS suppliersB
WHERE  suppliersA.city = suppliersB.city;
```

snum		snum
-----	+	-----
S1		S4
S1		S1
S2		S3
S2		S2
S3		S3
S3		S2
S4		S4
S4		S1
S5		S5
S6		S6

Queries in SQL

Suppliers database

Query: Retrieve all pairs of suppliers that are located in the same city.

```
SELECT DISTINCT suppliersA.snum, suppliersB.snum
FROM   suppliers AS suppliersA, suppliers AS suppliersB
WHERE  suppliersA.city = suppliersB.city;
```

snum		snum
-----	+	-----
S1		S4
S1		S1
S2		S3
S2		S2
S3		S3
S3		S2
S4		S4
S4		S1
S5		S5
S6		S6

Same result!

Queries in SQL

Suppliers database

Query: Retrieve all pairs of suppliers that are located in the same city.

```
SELECT suppliersA.snum, suppliersB.snum
FROM   suppliers AS suppliersA, suppliers AS suppliersB
WHERE  suppliersA.city = suppliersB.city
AND    suppliersA.snum != suppliersB.snum;
```

snum		snum
S1		S4
S2		S3
S3		S2
S4		S1

Better result

Queries in SQL

Suppliers database

Query: Retrieve all pairs of suppliers that are located in the same city.

```
SELECT suppliersA.snum, suppliersB.snum
FROM   suppliers AS suppliersA, suppliers AS suppliersB
WHERE  suppliersA.city = suppliersB.city
AND    suppliersA.snum < suppliersB.snum;
```

snum		snum
S1		S4
S2		S3

Best result

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all suppliers who supply both a red part and a blue part.

```
select  suppliers.snum, suppliers.sname
from    suppliers, shipments as shipmentsRed, shipments as shipmentsBlue,
        parts as partsRed, parts as partsBlue
where   partsBlue.color = 'Blue'
and     partsRed.color = 'Red'
and     suppliers.snum = shipmentsRed.snum
and     suppliers.snum = shipmentsBlue.snum
and     partsRed.pnum = shipmentsRed.pnum
and     partsBlue.pnum = shipmentsBlue.pnum;
```

snum		sname
S3		Blake

(1 row)

Queries in SQL

Suppliers database

Query: Retrieve the name, number, and color of all parts that are supplied by supplier S3.

```
SELECT DISTINCT parts.pname, parts.pnum, parts.color
FROM   parts JOIN shipments USING (pnum)
WHERE  shipments.snum = 'S3';
```

pname	pnum	color
Screw	P3	Blue
Screw	P4	Red

(2 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name, number, and color of all parts that are **NOT** supplied by supplier S3.

```
SELECT DISTINCT parts.pname, parts.pnum, parts.color
FROM   parts JOIN shipments USING (pnum)
WHERE  shipments.snum != 'S3';
```

pname	pnum	color
Screw	P4	Red
Screw	P3	Blue
Cog	P6	Red
Bolt	P2	Green
Nut	P1	Red
Cam	P5	Blue

(6 rows)

Incorrect. This returns the parts shipped by suppliers other than S3.

Queries in SQL

Suppliers database

Query: Retrieve the name, number, and color of all parts that are **NOT** supplied by supplier S3.

```
SELECT parts.pname, parts.pnum, parts.color
FROM   parts
EXCEPT
SELECT parts.pname, parts.pnum, parts.color
FROM   parts JOIN shipments USING (pnum)
WHERE  shipments.snum = 'S3';
```

pname	pnum	color
Gear	P7	Yellow
Cog	P6	Red
Bolt	P2	Green
Nut	P1	Red
Cam	P5	Blue

(5 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of suppliers who are currently not shipping any parts.

```
(SELECT sname, snum
FROM   suppliers)
EXCEPT
(SELECT sname, snum
FROM   suppliers JOIN shipments USING (snum));
```

sname	snum
Brown	S6

(1 row)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of suppliers who are currently not shipping any parts.

Start with a left join:

```
SELECT suppliers.sname, shipments.snum
FROM   suppliers LEFT JOIN shipments USING (snum);
```

sname		snum
-----+-----		
Smith		S1
Smith		S1
Jones		S2
Jones		S2
Jones		S2
Jones		S2
Jones		S2
Jones		S2
Jones		S2
Jones		S2
Jones		S2
. . .		

. . .
Blake S3
Blake S3
Clark S4
Clark S4
Adams S5
Adams S5
Adams S5
Adams S5
Smith S1
Smith S1
Brown
(21 rows)

← **Brown doesn't ship any parts.**

Queries in SQL

Suppliers database

Query: Retrieve the name and number of suppliers who are currently not shipping any parts.

```
SELECT suppliers.sname, shipments.snum  
FROM   suppliers LEFT JOIN shipments USING (snum)  
WHERE  shipments.snum IS NULL;
```

sname	snum
Brown	

(1 row)

← But we want their supplier number to show up.

Queries in SQL

Suppliers database

Query: Retrieve the name and number of suppliers who are currently not shipping any parts.

```
SELECT suppliers.sname, suppliers.snum
FROM   suppliers LEFT JOIN shipments USING (snum)
WHERE  shipments.snum IS NULL;
```

sname	snum
Brown	S6

(1 row)

Comparing to NULL

[more to come...]

NULL is considered a special marker that indicates the absence of value.

NULL corresponds to Codd's original idea of "missing or inapplicable data."

NULL is not a member of any data domain and is not considered a "value" at all. (So something can't "equal" NULL.)

```
SELECT suppliers.sname, suppliers.snum
FROM   suppliers LEFT JOIN shipments USING (snum)
WHERE  shipments.snum == NULL;
```

```
sname | snum
ERROR: operator does not exist: text == unknown
LINE 3: WHERE  shipments.snum == NULL;
                        ^
```

HINT: No operator matches the given name and argument type(s). You might need to add explicit type casts.-----+-----

```
SELECT suppliers.sname, suppliers.snum
FROM   suppliers LEFT JOIN shipments USING (snum)
WHERE  shipments.snum IS NULL;
```

```
sname | snum
-----+-----
Brown | S6
(1 row)
```


Aggregate functions

Queries in SQL

Suppliers database

Query: Retrieve the total number of suppliers in the database.

```
SELECT COUNT(snum)
FROM   suppliers;
```

```
count
-----
      6
(1 row)
```

```
SELECT COUNT(*)
FROM   suppliers;
```

```
count
-----
      6
(1 row)
```

```
SELECT COUNT(DISTINCT snum)
FROM   suppliers;
```

```
count
-----
      6
(1 row)
```

Queries in SQL

Suppliers database

Query: Retrieve the total number of suppliers currently shipping parts.

```
SELECT COUNT (snum)
FROM  shipments;
```

Incorrect

```
count
-----
      20
(1 row)
```

```
SELECT COUNT (DISTINCT snum)
FROM  shipments;
```

```
count
-----
       5
(1 row)
```

Queries in SQL

Suppliers database

Query: Retrieve the total number of shipments of part P2.

```
SELECT COUNT(*)  
FROM   shipments  
WHERE  pnum = 'P2';
```

```
count  
-----  
      2  
(1 row)
```

Queries in SQL

Suppliers database

Query: Retrieve the total quantity of part P2 that is being shipped.

```
SELECT SUM(qty)
FROM   shipments
WHERE  pnum = 'P2';
```

```
sum
-----
300
(1 row)
```

Queries in SQL

A given DBMS will likely support many aggregate functions, but the standard set are:

AVG

COUNT

MIN

MAX

SUM

Group by and Having

Queries in SQL

Suppliers database

Query: For each part being shipped, retrieve the total quantity being shipped.

```
SELECT pnum, SUM(qty)
FROM   shipments
GROUP BY pnum;
```

pnum	sum
P2	300
P6	800
P4	600
P3	3300
P5	700
P1	900

(6 rows)

Thinking about GROUP BY conceptually ...

We can *conceptually* think of the GROUP BY clause arranging the table into the specified groups and then returning one result row per group.

Shipments

snum	pnum	jnum	qty
S1	P1	J1	200
S1	P1	J4	700
S2	P3	J1	400
S2	P3	J2	200
S2	P3	J3	200
S2	P3	J4	500
S2	P3	J5	600
S2	P3	J6	400
S2	P3	J7	800
S2	P5	J2	100
S3	P3	J1	200
S3	P4	J2	500
S4	P6	J3	300
S4	P6	J7	300
S5	P2	J2	200
S5	P2	J4	100
S5	P5	J5	500
S5	P5	J7	100
S1	P4	J1	100
S1	P6	J2	200

Shipments

snum	pnum	jnum	qty
S1	P1	J4	700
<u>S1</u>	<u>P1</u>	<u>J1</u>	<u>200</u>
S5	P2	J4	100
<u>S5</u>	<u>P2</u>	<u>J2</u>	<u>200</u>
S2	P3	J4	500
S2	P3	J5	600
S2	P3	J6	400
S2	P3	J7	800
S3	P3	J1	200
S2	P3	J1	400
S2	P3	J2	200
<u>S2</u>	<u>P3</u>	<u>J3</u>	<u>200</u>
S1	P4	J1	100
<u>S3</u>	<u>P4</u>	<u>J2</u>	<u>500</u>
S5	P5	J7	100
S2	P5	J2	100
<u>S5</u>	<u>P5</u>	<u>J5</u>	<u>500</u>
S1	P6	J2	200
S4	P6	J7	300
<u>S4</u>	<u>P6</u>	<u>J3</u>	<u>300</u>

Thinking about GROUP BY conceptually ...

```
SELECT pnum, SUM(qty)
FROM   shipments
GROUP BY pnum;
```

Shipments

snum	pnum	jnum	qty
S1	P1	J1	200
S1	P1	J4	700
S2	P3	J1	400
S2	P3	J2	200
S2	P3	J3	200
S2	P3	J4	500
S2	P3	J5	600
S2	P3	J6	400
S2	P3	J7	800
S2	P5	J2	100
S3	P3	J1	200
S3	P4	J2	500
S4	P6	J3	300
S4	P6	J7	300
S5	P2	J2	200
S5	P2	J4	100
S5	P5	J5	500
S5	P5	J7	100
S1	P4	J1	100
S1	P6	J2	200

Shipments

snum	pnum	jnum	qty
S1	P1	J4	700
<u>S1</u>	<u>P1</u>	<u>J1</u>	<u>200</u>
S5	P2	J4	100
<u>S5</u>	<u>P2</u>	<u>J2</u>	<u>200</u>
S2	P3	J4	500
S2	P3	J5	600
S2	P3	J6	400
S2	P3	J7	800
S3	P3	J1	200
S2	P3	J1	400
S2	P3	J2	200
<u>S2</u>	<u>P3</u>	<u>J3</u>	<u>200</u>
S1	P4	J1	100
<u>S3</u>	<u>P4</u>	<u>J2</u>	<u>500</u>
S5	P5	J7	100
S2	P5	J2	100
<u>S5</u>	<u>P5</u>	<u>J5</u>	<u>500</u>
S1	P6	J2	200
S4	P6	J7	300
<u>S4</u>	<u>P6</u>	<u>J3</u>	<u>300</u>

Quick digression – order of evaluation

SQL is more declarative than procedural, but there is an order of evaluation that is usually observed.

Different DBMS implementations could behave differently, but this is typical.

```
5 ---SELECT <attribute list>  
1 ---FROM   <table list>  
2 ---[WHERE <condition>]  
3 ---[GROUP BY <grouping attribute(s)>  
4 ---[HAVING <group condition>]]  
6 ---[ORDER BY <attribute list>]
```

Queries in SQL

Suppliers database

Query: For each part in the database, retrieve the total quantity being shipped.

```
SELECT parts.pnum, SUM(qty)
FROM   parts LEFT JOIN shipments USING (pnum)
GROUP BY pnum;
```

pnum	sum
P4	600
P3	3300
P1	900
P2	300
P6	800
P7	
P5	700

(7 rows)

NULL

Not zero! [More later ...]

Queries in SQL

Suppliers database

Query: For each supplier currently shipping at least one part, list their supplier number, status, and total quantity of parts being shipped. List the result in descending order of status.

```
SELECT  snum, status, SUM(qty) AS totalQty
FROM    suppliers JOIN shipments USING (snum)
GROUP BY snum
ORDER BY status DESC;
```

snum	status	totalQty
S5	30	900
S3	30	700
S1	20	1200
S4	20	600
S2	10	3200

(5 rows)

Queries in SQL

Suppliers database

Query: Retrieve the part numbers for all parts being shipped by more than one supplier.

```
SELECT pnum
FROM   shipments
GROUP BY pnum
HAVING COUNT(DISTINCT snum) > 1;
```

```
pnum
-----
P3
P4
P5
P6
(4 rows)
```

Queries in SQL

Suppliers database

Query: For those suppliers shipping more than one kind of part, retrieve their supplier number.

```
SELECT snum
FROM   shipments
GROUP BY snum
HAVING COUNT(DISTINCT pnum) > 1;
```

```
snum
-----
S1
S2
S3
S5
(4 rows)
```

Queries in SQL

Suppliers database

Query: For those suppliers shipping more than one kind of part, retrieve their supplier number **and all associated part numbers**.

Incorrect

```
SELECT snum, pnum
FROM   shipments
GROUP BY snum
HAVING COUNT(DISTINCT pnum) > 1;
```

ERROR: column "shipments.pnum" must appear in the GROUP BY clause or be used in an aggregate function

snum		pnum
S1		P1
S1		P4
S1		P6
S2		P3
S2		P5
S3		P3
S3		P4
S5		P2
S5		P5

(9 rows)

Queries in SQL

Suppliers database

Query: For those suppliers shipping more than one kind of part, retrieve their supplier number **and all associated part numbers**.

Incorrect

```
SELECT snum, pnum
FROM   shipments
GROUP BY snum, pnum
HAVING COUNT(DISTINCT pnum) > 1;
```

```
snum | pnum
-----+-----
(0 rows)
```

```
snum | pnum
-----+-----
S1   | P1
S1   | P4
S1   | P6
S2   | P3
S2   | P5
S3   | P3
S3   | P4
S5   | P2
S5   | P5
(9 rows)
```

Subqueries

Queries in SQL

Suppliers database

Part of the data source
needs to be computed

Query: For those suppliers shipping more than one kind of part, retrieve their supplier number **and all associated part numbers**.

```
SELECT DISTINCT shipments.snum, shipments.pnum
FROM   shipments
      JOIN (SELECT snum
            FROM   shipments
            GROUP BY snum
            HAVING COUNT(DISTINCT pnum) > 1) AS mult_supplier
      USING (snum)
ORDER BY snum, pnum;
```

snum	pnum
S1	P1
S1	P4
S1	P6
S2	P3
S2	P5
S3	P3
S3	P4
S5	P2
S5	P5

(9 rows)

Subqueries in SQL

Subqueries can be used in many places where the type of value that they return is appropriate.

The most common places to use a subquery are the FROM and WHERE clauses.

Operators for subqueries in the WHERE clause:

x **IN** (subquery) checks if x is in the result of the subquery

x **NOT IN** (subquery) checks if x is not in the result of the subquery


EXISTS (subquery) checks if the result of subquery is non-empty

NOT EXISTS (subquery) checks if the result of subquery is empty

x comparison-operator (subquery)

x comparison-operator **ALL** (subquery)

x comparison-operator **SOME** (subquery)

- 
- comparison-operator: =, !=, <, >, <=, >=
 - result of the subquery must have only one column
 - If only one tuple results from the subquery, no quantifier is needed
 - ALL: comparison true for all tuples in result
 - SOME: comparison true for some tuple in result

Queries in SQL

Suppliers database

Query: For those suppliers shipping more than one kind of part, retrieve their supplier number **and all associated part numbers**.

```
SELECT DISTINCT snum, pnum
FROM   shipments
WHERE  snum IN
        (SELECT snum
         FROM   shipments
         GROUP BY snum
         HAVING COUNT(DISTINCT pnum) > 1)
ORDER BY snum, pnum;
```

snum	pnum
S1	P1
S1	P4
S1	P6
S2	P3
S2	P5
S3	P3
S3	P4
S5	P2
S5	P5

(9 rows)

Queries in SQL

Suppliers database

Query: For those suppliers shipping more than one kind of part, retrieve their supplier number **and all associated part numbers**.

```
SELECT DISTINCT shipmentsA.snum, shipmentsA.pnum
FROM   shipments shipmentsA
WHERE  EXISTS
        (SELECT shipmentsB.snum
         FROM   shipments shipmentsB
         WHERE  shipmentsA.snum = shipmentsB.snum
         GROUP BY snum
         HAVING COUNT(DISTINCT pnum) > 1)
ORDER BY snum, pnum;
```

snum		pnum
S1		P1
S1		P4
S1		P6
S2		P3
S2		P5
S3		P3
S3		P4
S5		P2
S5		P5

(9 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of suppliers with current maximum status.

```
SELECT sname, snum
FROM   suppliers
WHERE  status = (SELECT MAX(status)
                  FROM   suppliers);
```

sname		snum
-----+-----		
Blake		S3
Adams		S5

(2 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of suppliers who are currently not shipping any parts.

```
SELECT sname, snum
FROM   suppliers
WHERE  snum NOT IN
      (SELECT DISTINCT snum
       FROM   shipments);
```

sname		snum
-----+-----		
Brown		S6

Subqueries in SQL

Retrieve the name and number of suppliers who are currently shipping at least one part.

```
SELECT sname, snum
FROM   suppliers
WHERE  snum IN
      (SELECT DISTINCT snum
       FROM   shipments);
```

```
SELECT sname, snum
FROM   suppliers
WHERE  snum = SOME
      (SELECT DISTINCT snum
       FROM   shipments);
```

```
SELECT sname, snum
FROM   suppliers
WHERE  EXISTS
      (SELECT DISTINCT snum
       FROM   shipments
       WHERE shipments.snum = suppliers.snum);
```

Subqueries in SQL

Retrieve the name and number of suppliers who are currently not shipping any parts.

```
SELECT sname, snum
FROM   suppliers
WHERE  snum NOT IN
      (SELECT DISTINCT snum
       FROM   shipments);
```

```
SELECT sname, snum
FROM   suppliers
WHERE  snum != ALL
      (SELECT DISTINCT snum
       FROM   shipments);
```

```
SELECT sname, snum
FROM   suppliers
WHERE  NOT EXISTS
      (SELECT DISTINCT snum
       FROM   shipments
       WHERE shipments.snum = suppliers.snum);
```

Subqueries in SQL

Retrieve the name and number of suppliers who have the current maximum status.

```
SELECT sname, snum
FROM   suppliers
WHERE  status >= ALL
      (SELECT status
       FROM   suppliers);
```

Retrieve the name and number of suppliers who don't have the minimum status.

```
SELECT sname, snum
FROM   suppliers
WHERE  status > SOME
      (SELECT status
       FROM   suppliers);
```

Queries in SQL

Suppliers database

Revisit using subqueries.

Query: Retrieve the name and number of all suppliers who supply **either** a red part **or** a blue part.

```
SELECT DISTINCT suppliers.sname, suppliers.snum
FROM   suppliers, shipments, parts
WHERE  suppliers.snum = shipments.snum
AND    parts.pnum = shipments.pnum
AND   parts.color IN ('Red', 'Blue');
```

sname		snum
-----+-----		
Jones		S2
Blake		S3
Smith		S1
Clark		S4
Adams		S5

(5 rows)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all suppliers who supply **both** a red part **and** a blue part.

```
SELECT sname, snum
FROM   suppliers
WHERE  snum IN
        (SELECT sup.snum
         FROM   suppliers AS sup JOIN shipments USING (snum)
         JOIN parts USING (pnum)
         WHERE  parts.color = 'Red')
AND    snum IN
        (SELECT sup.snum
         FROM   suppliers AS sup JOIN shipments USING (snum)
         JOIN parts USING (pnum)
         WHERE  parts.color = 'Blue');
```

sname	snum
Blake	S3

(1 row)

Queries in SQL

Suppliers database

Query: Retrieve the suppliers who are currently shipping to all the projects.

Two common approaches: (1) Implement using only core relational algebra operators. (2) Apply a logical quantification tautology.

Queries in SQL

Suppliers database

Query: Retrieve the name and number of suppliers who are currently shipping all the parts that are carried.

```
select distinct snum from shipments
except
select snum
  from (select snum, pnum
        from (select snum from shipments) as t1,
              (select pnum from parts) as t2
        except
        select snum, pnum from shipments) as t3;
```

```
sname | snum
-----+-----
(0 rows)
```

Queries in SQL

Suppliers database

Query: Retrieve the name and number of suppliers who are currently shipping all the red parts that are being shipped.

```
select distinct snum from shipments
except
select snum
  from (select snum, pnum
        from (select snum from shipments) as t1,
              (select pnum from parts where color = 'Red') as t2
        except
        select snum, pnum from shipments) as t3;
```

sname	snum
Smith	S1

(1 row)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of suppliers who are currently shipping all the red parts that are being shipped.

```
SELECT sname, snum
FROM   suppliers
WHERE  NOT EXISTS
      (SELECT *
       FROM   parts
       WHERE  color = 'Red'
       AND    NOT EXISTS
            (SELECT *
             FROM   shipments
             WHERE  shipments.pnum = parts.pnum
             AND    shipments.snum = suppliers.snum));
```

sname		snum
-----+-----		
Smith		S1

(1 row)

Queries in SQL

Suppliers database

Query: Retrieve the name and number of the parts that are being shipped to all the projects in Paris.

```
SELECT pname, pnum
FROM parts
WHERE NOT EXISTS
    (SELECT *
     FROM projects
     WHERE city = 'Paris'
     AND NOT EXISTS
        (SELECT *
         FROM shipments
         WHERE shipments.jnum = projects.jnum
         AND shipments.pnum = parts.pnum));
```

pname	pnum
Nut	P1
Screw	P3
Screw	P4

(3 rows)

Things that can go in the attribute list

SQL SELECT statement – basic examples

```
SELECT *  
FROM   suppliers  
ORDER BY status;
```

snum	sname	status	city
S2	Jones	10	Paris
S6	Brown	15	Berlin
S1	Smith	20	London
S4	Clark	20	London
S3	Blake	30	Paris
S5	Adams	30	Athens

(6 rows)

SQL SELECT statement – basic examples

```
SELECT pnum, weight * 454 AS "Weight in Grams"  
FROM   parts;
```

pnum	Weight in Grams
P1	5448
P2	7718
P3	7718
P4	6356
P5	5448
P6	8626
P7	8172

(7 rows)

Views

Queries in SQL

Suppliers database

Query: Retrieve the name and number of all parts that are carried in any of the following colors: red, yellow, green.

```
CREATE VIEW Red_Parts AS
```

```
(SELECT pname, pnum FROM parts WHERE color = 'Red');
```

```
CREATE VIEW Yellow_Parts AS
```

```
(SELECT pname, pnum FROM parts WHERE color = 'Yellow');
```

```
CREATE VIEW Green_Parts AS
```

```
(SELECT pname, pnum FROM parts where color = 'Green';
```

```
(SELECT * FROM Red_Parts) UNION (SELECT * FROM Green_Parts) UNION  
(SELECT * FROM Yellow_Parts);
```

pname	pnum
Nut	P1
Gear	P7
Bolt	P2
Cog	P6
Screw	P4

(5 rows)

Conditional functions

Queries in SQL

Suppliers database

Query 17 For each part in the database, retrieve the total quantity being shipped.

```
SELECT parts.pnum, COALESCE(SUM(qty), 0)
FROM   parts LEFT JOIN shipments USING (pnum)
GROUP BY pnum;
```

pnum	coalesce
P4	600
P3	3300
P1	900
P2	300
P6	800
P7	0
P5	700

(7 rows)

The COALESCE function returns the first of its arguments that is not NULL.

Summary: Some tips before you leave

Test and Troubleshoot

- DO NOT wait until the end to test queries
- Test after each join or filter
- Are you getting the results you expect?
- Start small and go step-by-step when troubleshooting a query

Format and Comment

- Use correct formatting
- Comment strategically
- Clean code and comments when you revisit and hand off code

Review

- Always review old queries
- Business rules
- Data changes, data indicators
- Work the problem from beginning to end

Haven't had enough SQL? Or do you want to keep practicing to improve your skills? SQL puzzles are a great way to do this! Below is a very popular resources for practicing SQL Puzzles.

[SQL Authority: SQL Puzzles](#)

In addition, many of you may desire to get a new job or position in this are. Below is a resource that includes quizzes and is recommended by many recruiters to practice SQL for a data science interview.

[SQLZOO](#)